

# "Cortex AI"

**Tagline:** *Your Development Team's AI Brain - Learns, Decides, Automates*

## The One-Liner Pitch:

"Cortex AI monitors your entire development workflow, learns from your team's patterns using reinforcement learning, and autonomously makes decisions to optimize your pipeline - from code review to deployment."

---

## Why This WINS

### 1. Strong Innovation Story

- Not just automation - **autonomous decision-making**
- Not just monitoring - **predictive intelligence**
- Not just AI summaries - **RL-trained on YOUR team's data**

### 2. Clear Value Proposition

Every developer immediately understands: "An AI that learns my team's patterns and automates decisions? That's game-changing."

### 3. Perfect Tool Integration (All 5 tools are ESSENTIAL, not forced)

### 4. Demonstrable Impact

Live demo shows real decisions being made in real-time

### 5. Scalable Vision

Judges can immediately see how this becomes a real product

---

## THE CORE CONCEPT

### What Cortex AI Does:

#### Think of it as an AI team member that:

1. **Observes** - Monitors your GitHub repo, PRs, CI/CD, deployments
2. **Learns** - Uses RL to understand what "good" looks like for your team
3. **Summarizes** - Provides intelligent insights across your workflow
4. **Decides** - Makes autonomous decisions based on learned patterns
5. **Acts** - Executes decisions through automated workflows

## The Magic Moment (What Makes Judges Go "Wow"):

A developer pushes code → Cortex AI analyses it → Learns similar to past successful patterns → Decides to fast-track it → Auto-approves review → Triggers deployment → Monitors success → Updates its decision model.

**All autonomous. All learned. All demonstrated live.**

---

## HOW EACH TOOL CREATES THE SYSTEM

### 1. Cline CLI - The Command Interface

**Role:** Direct interaction layer for developers

**Commands You'll Build:**

```
cline analyze      # Analyze current repo state  
cline predict     # Predict PR merge time, deployment risk  
cline decide [action] # Ask Cortex to make a decision  
cline deploy [env]   # Intelligent deployment with AI checks  
cline learn        # Trigger RL training update  
cline insights      # Get AI-generated insights  
cline status        # Full system status
```

**Why It's Essential:** Gives developers direct access to AI brain. Showcases automation.

---

### 2. Kestra - The Orchestration Engine

**Role:** Workflow automation and data pipeline backbone

**Workflows You'll Build:**

#### Workflow 1: Data Collection Pipeline

- Continuously pulls data from GitHub (commits, PRs, reviews, CI results)
- Stores in structured format for AI training
- Runs every 15 minutes

#### Workflow 2: AI Decision Engine

- Triggered on PR creation/update
- Collects contextual data (author history, file changes, test coverage)
- Feeds to Kestra AI Agent
- Executes decisions

### **Workflow 3: Deployment Orchestration**

- Handles staging → production pipeline
- AI-driven deployment decisions
- Rollback triggers based on metrics

**Why It's Essential:** The nervous system connecting everything. Shows real orchestration complexity.

---

### **3. Kestra AI Agent - The Decision Maker**

#### **BONUS POINTS**

**Role:** Summarizes data AND makes autonomous decisions

**Decision Examples:**

#### **Decision 1: PR Review Priority**

- Summarizes: PR complexity, author reliability, urgency signals
- Decides: "High priority - assign senior reviewer immediately" vs "Low risk - auto-approve after CI"

#### **Decision 2: Deployment Go/No-Go**

- Summarizes: Test results, code coverage delta, similar past deployments
- Decides: "Safe to deploy" vs "Hold - similar pattern caused rollback 3 weeks ago"

#### **Decision 3: Resource Allocation**

- Summarizes: CI queue length, historical patterns
- Decides: "Allocate extra runners - high traffic period predicted"

**Why It's Essential:** This gets you BONUS POINTS. Shows AI not just summarizing but actually making calls.

---

### **4. Oumi - The Learning Brain**

**Role:** RL fine-tuning for personalized intelligence

**What You'll Train:**

**Dataset:** Your team's historical data

- Successful PRs vs problematic ones
- Fast deployments vs failed ones
- Good code patterns vs bug-prone patterns

**RL Training Goal:** Learn to predict outcomes and recommend actions that maximize:

- Code quality

- Deployment success rate
- Developer velocity

#### The Model Learns:

- "PRs from developer X touching auth code need extra review"
- "Deployments on Friday afternoons have 3x higher rollback rate"
- "Small PRs with high test coverage can be fast-tracked"

**Why It's Essential:** This is your INNOVATION showcase. Shows real ML, not just API calls.

---

## 5. Vercel - The Intelligence Dashboard

**Role:** Real-time visualization and interaction interface

#### Dashboard Sections:

##### 1. Command Center

- Live workflow status
- Recent AI decisions with reasoning
- System health metrics

##### 2. Intelligence Feed

- Real-time insights stream
- Decision explanations
- Prediction confidence scores

##### 3. Learning Dashboard

- RL training progress
- Model accuracy metrics
- Decision success rate over time

##### 4. Interactive Console

- Chat interface to query Cortex
- Manual decision overrides
- Training data exploration

## 5. Analytics

- Team velocity trends
- Deployment success rates
- Time saved through automation

**Why It's Essential:** Makes everything visible. Perfect for demo video. Shows it's a real product.

---

## 6. CodeRabbit - The Quality Guardian

**Role:** Code review throughout your development + training data source

**How You'll Use It:**

**During Development (Visible Activity):**

- Every feature gets a PR with CodeRabbit review
- Document CodeRabbit suggestions you implement
- Show 10-15 PRs with clear CodeRabbit activity

**As Data Source:**

- Use CodeRabbit's reviews as training signal for Oumi
- "CodeRabbit flagged security issue" → "Cortex learns to catch similar patterns"
- Makes the system smarter over time

**Why It's Essential:** Shows development best practices + feeds your AI training.

---

## THE DEMO NARRATIVE (5-7 Minutes)

### Act 1: The Problem (1 min)

"Development teams waste 40% of time on decisions that could be automated. Let me show you Cortex AI - an AI that learns your team's patterns and makes decisions for you."

### Act 2: The Learning (1.5 min)

- Show Oumi RL training on sample repository data
- Explain: "Cortex learned from 200 PRs, 50 deployments"
- Display model accuracy metrics improving

### Act 3: The Intelligence (2 min)

**Live Demo:**

1. Create a PR in demo repo
2. Show Kestra workflow trigger
3. Kestra AI Agent analyzes and decides
4. Dashboard updates with decision reasoning
5. Run cline predict - see AI predictions
6. Show decision execution

#### **Act 4: The Impact (1.5 min)**

- Show analytics: "Cortex saved 15 hours this week"
- Display decision accuracy: 92%
- Explain scalability: "Learns continuously, gets smarter over time"

#### **Act 5: The Vision (1 min)**

"Today: PR reviews and deployments. Tomorrow: Sprint planning, incident response, resource optimization. Cortex becomes your team's AI brain."

---

### **4-DAY EXECUTION PLAN**

#### **DAY 1: Foundation + Quick Win**

- Repository setup with proper structure
- Vercel deployment (basic Next.js app)
- CodeRabbit integration enabled
- Kestra Docker setup
- First PR with CodeRabbit review (setup docs)
- Basic dashboard skeleton deployed
- Cline CLI installation and first 2 commands
- First Kestra workflow (GitHub data collection)
- Basic data pipeline working
- 2-3 PRs with CodeRabbit reviews

**Deliverable:** All tools connected, basic data flowing, deployed dashboard

---

#### **DAY 2: Intelligence Layer**

- Oumi environment setup
- Prepare training dataset (can use synthetic + some real data)
- Start RL fine-tuning (let it run while you work on other things)
- Build 3 more Cline commands
- 3-4 PRs with CodeRabbit
- Kestra AI Agent integration
- Build first decision workflow (PR analysis)
- Test decision-making logic

- Dashboard: Add Intelligence Feed section
- 2-3 PRs

**Deliverable:** AI Agent making first decisions, RL training running, CLI functional

---

### **DAY 3: Integration + Polish**

- Complete Oumi RL training
- Integrate trained model into decision pipeline
- Build deployment decision workflow
- Complete remaining Cline commands
- Dashboard: Add Learning Dashboard section
- 3-4 PRs
- End-to-end testing
- Dashboard: Complete all sections
- Build interactive console
- Polish UI/UX
- 2-3 PRs for bug fixes

**Deliverable:** Fully integrated system, impressive dashboard, all features working

---

### **DAY 4: Demo Prep + Submission**

- Create demo repository with realistic activity
- Record demo video (multiple takes)
- Write comprehensive README with:
  - Project overview
  - Architecture diagrams
  - Setup instructions
  - Each tool's role explained
  - Screenshots/GIFs
- Final testing
  - Polish demo video (editing, captions)
- Final Vercel deployment check
- Verify all CodeRabbit activity visible

- Create submission materials
- Final PR with complete docs
- Submit

**Deliverable:** Submitted project with killer demo

---

## SUCCESS METRICS

### Technical Excellence

- All 5 tools integrated seamlessly
- Working RL model with visible training
- AI making real decisions
- Clean, professional code

### Innovation

- Autonomous decision-making (rare in hackathons)
- RL personalization (technical depth)
- Holistic workflow intelligence (big thinking)

### Impact

- Solves real pain point
- Quantifiable benefits (time saved, accuracy)
- Clear scalability path

### Execution

- Polished dashboard
  - Professional demo video
  - Comprehensive documentation
  - Visible development process (CodeRabbit PRs)
- 

## RISK MITIGATION

### Risk 1: Oumi RL Training Takes Too Long

#### Mitigation:

- Use small dataset (200-500 examples)
- Train overnight Day 1 → Day 2
- Have backup: Use pre-trained model with light fine-tuning

## **Risk 2: Kestra AI Agent Complex**

### **Mitigation:**

- Start with simple summarization
- Add decision logic incrementally
- Worst case: 1-2 decisions are enough for demo

## **Risk 3: Integration Issues**

### **Mitigation:**

- Test integrations daily
- Build modularly (each component works standalone)
- Have backup demo data ready

## **Risk 4: Time Crunch**

### **Mitigation:**

- Built-in buffer (plan for 10 hours/day, can go 12 if needed)
  - Prioritize core features (marked  in plan)
  - Can cut: Some Cline commands, some dashboard sections
- 

## **DELIVERABLES CHECKLIST**

### **GitHub Repository**

- Clean project structure
- 15+ PRs with CodeRabbit reviews
- Comprehensive README
- Architecture documentation
- Setup instructions
- All source code

### **Vercel Deployment**

- Live dashboard URL
- Stable and responsive
- All features functional
- Professional UI

### **README**

- Project overview with "why"

- Architecture diagram
- Tool integration explanations
- Setup/installation guide
- Usage examples
- Screenshots/GIFs
- Demo video link
- Team information

#### **Demo Video**

- 5-7 minutes
  - Clear narrative
  - Live demonstrations
  - Impact metrics
  - Professional quality
  - Captions/annotations
- 

## **Tech Stack Breakdown - Cortex AI**

### **MANDATORY TOOLS (Competition Requirements)**

#### **1. Cline CLI**

- **Domain:** Command-line interface / Developer tools
- **Usage:** Build custom automation commands for the system
- **Tech:** Node.js/TypeScript based CLI tool

#### **2. Kestra**

- **Domain:** Backend orchestration / Workflow engine
- **Usage:** Automate workflows, data pipelines, AI agent integration
- **Tech:** YAML workflows, REST API, Docker deployment
- **AI Agent:** Built-in feature for data summarization + decision-making

#### **3. Oumi**

- **Domain:** Machine Learning / AI training
- **Usage:** Reinforcement Learning fine-tuning on team's code patterns
- **Tech:** Python library, PyTorch-based

#### 4. Vercel

- **Domain:** Frontend hosting / Deployment platform
- **Usage:** Deploy the dashboard web application
- **Tech:** Next.js deployment

#### 5. CodeRabbit

- **Domain:** Development process / Code review
  - **Usage:** PR reviews throughout development, training data source
  - **Tech:** GitHub integration
- 

### Supporting Technologies

#### Frontend (Dashboard)

- **Framework:** Next.js 14 (React)
- **Styling:** Tailwind CSS
- **UI Components:** shadcn/ui
- **Charts:** Recharts
- **Deployment:** Vercel
- **Language:** TypeScript

#### Backend/API

- **API Routes:** Next.js API routes
- **Database:** Optional - PostgreSQL/SQLite (for storing metrics)
- **OR:** Simple JSON file storage
- **Language:** TypeScript/Node.js

#### Orchestration Layer

- **Kestra:** Workflow orchestration
- **Docker:** Running Kestra instance
- **Language:** YAML (Kestra workflows) + Python (custom tasks)

#### ML/AI Layer

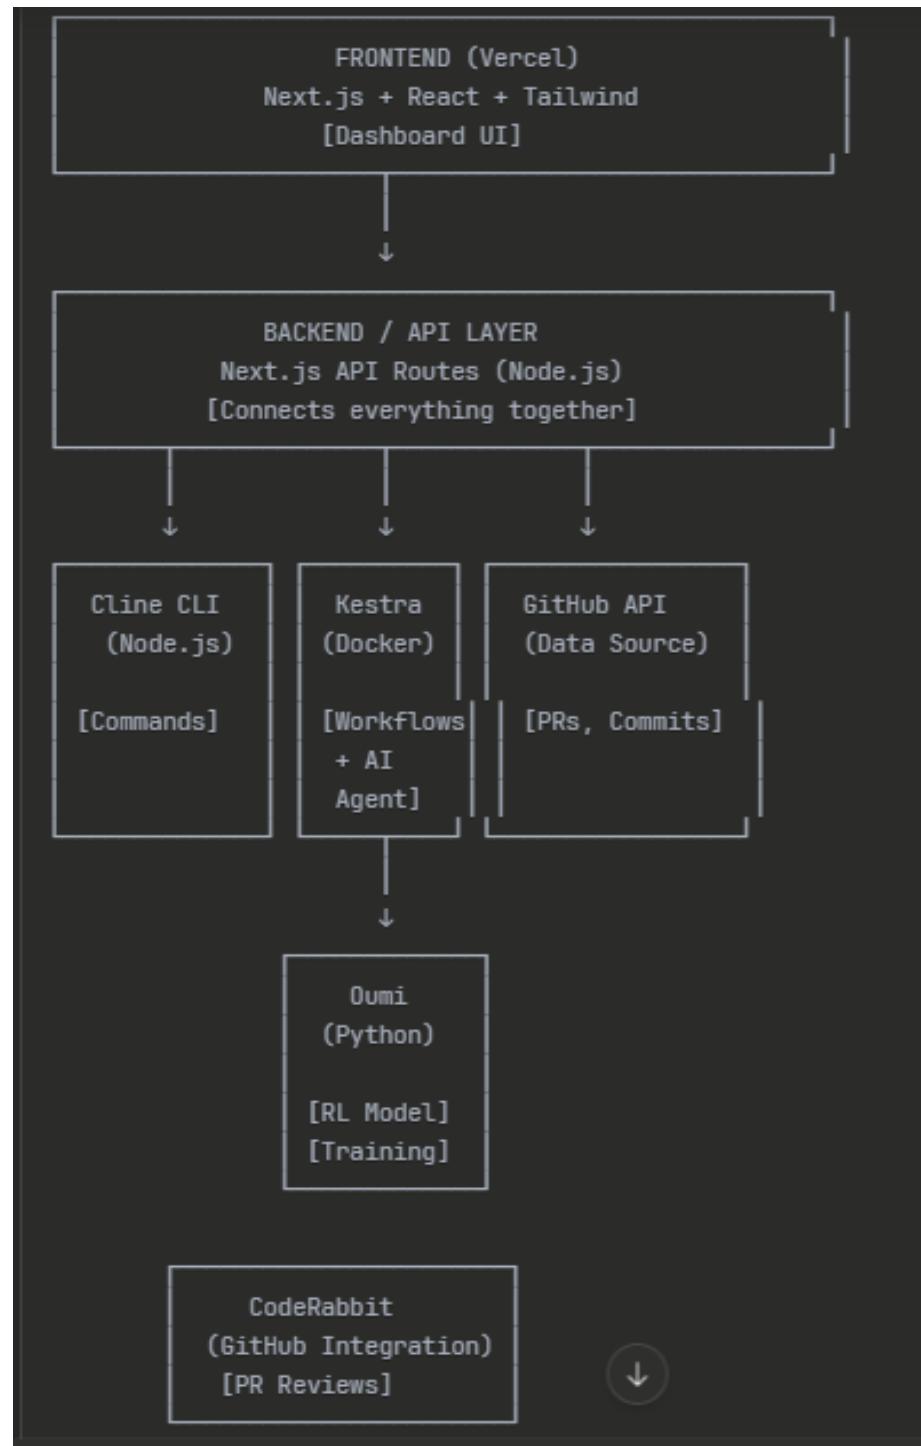
- **Oumi:** RL fine-tuning
- **Python:** Training scripts
- **PyTorch:** Under the hood (via Oumi)
- **Dataset:** JSON/CSV format

## CLI Tool

- **Cline CLI:** Command interface
- **Language:** Node.js/TypeScript
- **Integration:** Calls Kestra APIs, GitHub APIs

## Data Sources

- **GitHub API:** Pull repo data, PRs, commits
- **REST APIs:** Communication between components



Component	Tool/Tech	Domain	Language
Dashboard	Next.js + Vercel	Frontend	TypeScript
Styling	Tailwind CSS	Frontend	CSS
API	Next.js API Routes	Backend	TypeScript
CLI	Cline	Developer Tool	Node.js
Workflows	Kestra + AI Agent	Orchestration	YAML + Python
ML Training	Oumi	AI/ML	Python
Deployment	Vercel	Hosting	-
Code Review	CodeRabbit	Dev Process	-
Data Source	GitHub API	External	REST
Container	Docker	Infrastructure	-

## Efficient Task Division - Cortex AI

### TEAM STRUCTURE

#### Person A (You): Full-Stack + Integration Lead

- **Strengths:** Frontend, Backend, AI
- **Role:** System architect, integration specialist, dashboard builder

#### Person B (Aditya): AI/ML Specialist

- **Strengths:** AI models
- **Role:** ML training, AI agent logic, intelligent decision-making

### PERSON A (YOU) - Responsibilities

#### Infrastructure & Setup

- GitHub repository setup
- Vercel deployment pipeline
- CodeRabbit integration
- Docker setup for Kestra
- All environment configurations

## **Frontend Development**

- Next.js dashboard (complete UI)
- All dashboard sections:
  - Command Center
  - Intelligence Feed
  - Analytics
  - Interactive Console
- Real-time updates
- Charts and visualizations

## **Backend/API Development**

- Next.js API routes
- GitHub API integration
- Data collection and storage
- API endpoints for dashboard

## **Cline CLI**

- All CLI commands (7 commands)
- CLI → API communication
- Command testing and documentation

## **Kestra Workflows**

- Workflow YAML files
- Data pipeline orchestration
- GitHub data collection workflow
- Deployment workflow
- Integration with AI components

## **Integration Work**

- Connect Oumi model to system
- Connect Kestra AI Agent to workflows
- End-to-end system integration
- Testing all components together

## **Documentation & Demo**

- README file

- Demo video recording
- Architecture diagrams
- Setup instructions

### **CodeRabbit Management**

- Creating all PRs
  - Managing reviews
  - Ensuring visible activity
- 

### **PERSON B (TEAMMATE) - Responsibilities**

#### **Oumi RL Training (PRIMARY FOCUS)**

- Setup Oumi environment
- Prepare training dataset
- Design RL training strategy
- Run RL fine-tuning
- Model evaluation and optimization
- Export trained model
- Create inference code
- Document training process

#### **Kestra AI Agent Logic (SECONDARY FOCUS)**

- Design decision-making algorithms
- Create prompts for AI Agent
- Define decision rules and thresholds
- Test AI Agent responses
- Document AI Agent behavior

### **AI Decision Engine**

- PR analysis logic
- Deployment risk assessment
- Prediction algorithms
- Confidence scoring system

### **Training Data Preparation**

- Data cleaning and preprocessing

- Feature engineering
- Dataset formatting

## HANDOFF POINTS (Critical Communication)

### From Person B → Person A

1. **Day 2 Evening:**
  - Trained model file (.pth or equivalent)
  - Inference code/function
  - Input/output format specification
2. **Day 3 Morning:**
  - Kestra AI Agent prompt templates
  - Decision rules documentation
  - Expected behavior examples
3. **Day 3 Afternoon:**
  - Final model version
  - Performance metrics
  - Edge case handling

### From Person A → Person B

1. **Day 1 Evening:**
  - Data format from GitHub API
  - Sample data for training
  - Integration endpoint specs
2. **Day 2 Evening:**
  - Integration API ready
  - Testing environment
  - Kestra workflow structure
3. **Day 3 Morning:**
  - Complete system for AI integration
  - Testing framework
  - Dashboard for visualizing AI decisions

# Integration Workflow

- **Modular design** - Each component has clear inputs/outputs
- **API-based communication** - Everything talks via REST APIs
- **Person A handles ALL integration** - You're the integration layer

**Progressive Integration:**

**Day 2 Afternoon:**

- You build API endpoints expecting Person B's model (mock it first)
- Test with dummy AI responses

**Day 3 Morning (MAIN INTEGRATION):**

- Person B gives you: trained model file + inference code
- You integrate model into your API routes (2-3 hours)

**Day 3 Afternoon (FINAL INTEGRATION):**

- Person B gives you: Kestra AI Agent prompts/logic
- You add to Kestra workflows (1-2 hours)
- Full system testing

---

**Integration Complexity: LOW**

1. **Person B delivers 2 simple things:**
  - Python function: predict(data) → decision
  - YAML/prompts: Kestra AI Agent configuration
2. **You just call them:**
  - API route calls Python function
  - Kestra workflow uses AI Agent
3. **No complex dependencies** - Everything is REST/function calls

---

**Quick Answer:**

- **Where?** Your Next.js API routes
- **When?** Day 3 (morning + afternoon)
- **How long?** 3-4 hours total
- **Easy?** YES - just function calls and API endpoints